

```
\documentclass{article}
```

```
\usepackage[utf8]{inputenc}
```

```
\usepackage{enumerate}
```

```
\usepackage{listings}
```

```
\title{Lab 8}
```

```
\author{Jack Jiang}
```

```
\date{April 22, 2021}
```

```
\begin{document}
```

```
\maketitle
```

```
\section*{Tasks}
```

```
\begin{enumerate}
```

```
\item Setup Coding Environment -- Done
```

```
\item Review Provided Code
```

```
\begin{enumerate}[(a)]
```

\item In the vector structure, the data is distributed by distributing the data amongst the number of processors used. Data is an array of floats and holds values of type float in it when allocated. The data then can be distributed amongst the other processors in use and it is distributed evenly using halo exchanges with size of the overall ghost padding region (padding) and size for the local padded region (N_{padded}). Data will start being distributed at the starting index of the unpadded local vector (r_0) and be stored in the padded index of the vector structure.

\item The function `inject_unpadded_data_into_Vector(float* data, Vector* v)` is responsible for copying data into the unpadded portion of the local vector. It accomplishes this by calculating the padded index by adding the index of the unpadded data and adds the value of the padding of the vector to it and then for each index in the local vector, its data at the padded index is assigned the value of the data array at the unpadded index.

```
\item
```

```
\begin{enumerate}[i.]
```

```
\item mpirun -np 2 ./test\_linear\_function 21 3
```

\item After injecting the unpadded data, I expect v to have data corresponding to the inject array with a padding size of $\frac{k-1}{2}$, where K is the window size, as padding on each side of v's data array, numbered from 0.0 to N-1.

\item After applying mirror boundary conditions, the data I expect in v is to be, as stated in the function header, in a format such that say the data has a padding of 3, so its ---0123---. Then, after the mirror, it'll become 2100123321.

\item After sending/receiving the ghost regions, I expect the data in v to have switched values in the ghost regions as the processes send and receive their values to/amongst each other.

\item After applying the average function, I expect the data in v_avg to be the running average of the data computed so far in the execution of the program.

\end{enumerate}

\end{enumerate}

\item Implement Routines

\begin{enumerate}

\item Done (Code pushed to Git)

\end{enumerate}

\item Questions

\begin{enumerate}[(a)]

\item The inject/extract operations can be performed in one line by having an independent pointer pointing to the values you want to inject/extract and injecting/extracting with the pointer.

\item The output from each process for the average vector:

\begin{lstlisting}[language=bash]

Average data

0.3333333433

1.0000000000

2.0000000000

3.0000000000

4.0000000000

```
5.0000000000
6.0000000000
7.0000000000
8.0000000000
9.0000000000

10.0000000000
11.0000000000
12.0000000000
13.0000000000
14.0000000000
15.0000000000
16.0000000000
17.0000000000
18.0000000000
19.0000000000
19.6666660309
```

```
\end{lstlisting}
```

The values at the end of the global vector have a different pattern because the values being averaged at these points are the averages including the mirrored padding.

```
\item
```

I would expect for the 0th entry of the 0th process's unpadded array to be 0.0. The value being averaged is the 0th entry and the mirrored boundary conditions before the 0th entry.

I would expect for the 9th entry of the 0th process's unpadded array to be 9.0. The values to be averaged would be the first 9 values, including the 3 ghost padding values in the front.

The output from each process for the average vector:

```
\begin{lstlisting}[language=bash]
```

```
Average data
```

40.0000000000

41.0000000000

42.0000000000

43.0000000000

44.0000000000

45.0000000000

46.0000000000

47.0000000000

47.8571434021

48.4285697937

48.7142868042

10.0000000000

11.0000000000

12.0000000000

13.0000000000

14.0000000000

15.0000000000

16.0000000000

17.0000000000

18.0000000000

19.0000000000

20.0000000000

21.0000000000

22.0000000000

23.0000000000

24.0000000000

25.0000000000

26.0000000000

27.0000000000

28.0000000000

29.0000000000

30.0000000000

31.0000000000

32.0000000000

33.0000000000

34.0000000000

35.0000000000

36.0000000000

37.0000000000

38.0000000000

39.0000000000

1.2857142687

1.5714285374

2.1428570747

3.0000000000

4.0000000000

5.0000000000

6.0000000000

7.0000000000

8.0000000000

9.0000000000

\end{lstlisting}

\item

All the entries from the average data output have indeed changed as shown below:

\begin{lstlisting}

Average data

0.3333333433

1.666666269

4.666665077

9.666669846

16.666660309

25.666660309

36.666679382

49.666679382

64.666641235

81.666641235

100.666641235

121.666641235

144.6666717529

169.6666717529

196.6666717529

225.6666717529

256.666564941

289.666564941

324.666564941

361.666564941

387.000000000

\end{lstlisting}

The values being printed to output from each process after averaging is the running average of the quadratic calculation at each entry.

`\end{enumerate}`

`\item` I received no assistance on this assignment. I did assist Shaan Chudasama in the concepts and coding portion, helping him get an idea of how to approach the functions and what each should do.

`\end{enumerate}`

`\end{document}`